

```
In [1]: import pandas as pd
```

```
In [2]: df = pd.read_csv("labeled_data.csv")
```

```
In [3]: df.head()
```

Out[3]:

	Unnamed: 0	count	hate_speech	offensive_language	neither	class	tweet
0	0	3	0	0	3	2	!!! RT @mayasolovely: As a woman you shouldn't...
1	1	3	0	3	0	1	!!!! RT @mleew17: boy dats cold...tyga dwn ba...
2	2	3	0	3	0	1	!!!!!! RT @UrKindOfBrand Dawg!!!! RT @80sbaby...
3	3	3	0	2	1	1	!!!!!!! RT @C_G_Anderson: @viva_based she lo...
4	4	6	0	6	0	1	!!!!!!!!!!!! RT @ShenikaRoberts: The shit you...

```
In [4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 24783 entries, 0 to 24782
Data columns (total 7 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Unnamed: 0            24783 non-null  int64
1   count                 24783 non-null  int64
2   hate_speech           24783 non-null  int64
3   offensive_language    24783 non-null  int64
4   neither               24783 non-null  int64
5   class                 24783 non-null  int64
6   tweet                 24783 non-null  object
dtypes: int64(6), object(1)
memory usage: 1.3+ MB
```

```
In [5]: data = df.drop(columns = ["Unnamed: 0" , "count" , "hate_speech" , "offensive_language" , "neith
```

```
In [6]: data["class"] = data["class"].astype("category")
```

## PRE - PROCESSING

- TOKENIZATION
- LOWER CASING
- REMOVAL OF WHITESPACES
- REMOVAL OF STOP WORDS
- REMOVAL OF PUNCTUATION
- LEMMATIZATION

```
In [7]: ► import string
import re
from nltk.corpus import stopwords
import nltk
nltk.download('stopwords')
nltk.download('wordnet')
from nltk.corpus import wordnet as wn
from nltk.stem.wordnet import WordNetLemmatizer

def text_cleaning(text):
    text = text.lower()
    exclude = set(string.punctuation) - {"'"}
    text = ''.join(ch for ch in text if ch not in exclude)
    text = re.sub(r"^[^a-zA-Z0-9'\s]", "", text)
    return text

def stopword_removal(text):
    words = text.split()
    stop_words = stopwords.words('english')
    filtered_words = [word for word in words if word not in stop_words]
    text = ' '.join(filtered_words)
    return text

def tokenization(text):
    tokens = text.split()
    return tokens

def lemmatization(tokens):
    lemmatizer = WordNetLemmatizer()
    lemmatized_tokens = [lemmatizer.lemmatize(token) for token in tokens]
    return lemmatized_tokens
```

```
[nltk_data] Downloading package stopwords to
[nltk_data] C:\Users\chon2\AppData\Roaming\nltk_data...
[nltk_data] Package stopwords is already up-to-date!
[nltk_data] Downloading package wordnet to
[nltk_data] C:\Users\chon2\AppData\Roaming\nltk_data...
[nltk_data] Package wordnet is already up-to-date!
```

```
In [8]: ► def pre_proc(text):
        cleaned_text = text_cleaning(text)
        no_stopwords = stopword_removal(cleaned_text)
        tokens = tokenization(no_stopwords)
        lemmatized_tokens = lemmatization(tokens)

        return lemmatized_tokens
```

## PRE - PROCESSING THE TWEET COLUMN

```
In [9]: ► data["preproc_tweet"] = data["tweet"].apply(pre_proc)
```

```
In [10]: data
```

```
Out[10]:
```

	class	tweet	preproc_tweet
0	2	!!! RT @mayasolovely: As a woman you shouldn't...	[rt, mayasolovely, woman, complain, cleaning, ...
1	1	!!!! RT @mleew17: boy dats cold...tyga dwn ba...	[rt, mleew17, boy, dat, coldtyga, dwn, bad, cu...
2	1	!!!!!! RT @UrKindOfBrand Dawg!!!! RT @80sbaby...	[rt, urkindofbrand, dawg, rt, 80sbaby4life, ev...
3	1	!!!!!!! RT @C_G_Anderson: @viva_based she lo...	[rt, cganderson, vivabased, look, like, tranny]
4	1	!!!!!!!!!!!! RT @ShenikaRoberts: The shit you...	[rt, shenikaroberts, shit, hear, might, true, ...
...	...	...	...
24778	1	you's a muthaf***in lie &#8220;@LifeAsKing: @2...	[you's, muthafin, lie, 8220lifeasking, 20pearl...
24779	2	you've gone and broke the wrong heart baby, an...	[gone, broke, wrong, heart, baby, drove, redne...
24780	1	young buck wanna eat!!.. dat nigguh like I ain...	[young, buck, wanna, eat, dat, nigguh, like, a...
24781	1	youu got wild bitches tellin you lies	[youu, got, wild, bitch, tellin, lie]
24782	2	~~Ruffled   Ntac Eileen Dahlia - Beautiful col...	[ruffled, ntac, eileen, dahlia, beautiful, col...

24783 rows × 3 columns

```
In [11]: data['class'].unique()
```

```
Out[11]: [2, 1, 0]
Categories (3, int64): [0, 1, 2]
```

```
In [12]: X = data["preproc_tweet"]
y = data["class"]
```

## Word2Vect - CONTINUOUS BAG OF WORDS

### STEPS:

- TRAINING THE WORD2VECT - CBOW MODEL
- AVERAGE WORD VECTORS FOR A SENTENCE
- TRANSFORMING TRAINING AND TEST DATA TO AVERAGE WORD VECTORS
- ENCODE LABELS TO NUMERIC VALUES
- CLASSIFICATION USING LOGISTIC REGRESSION
- EVALUATING THE MODEL
  - ACCURACY
  - CONFUSION MATRIX
  - CLASSIFICATION REPORT

```
In [13]: from gensim.models import Word2Vec
from sklearn.metrics import accuracy_score
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import confusion_matrix, classification_report
import numpy as np

from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

w2v_model_cbow = Word2Vec(sentences=X_train, vector_size=100, window=5, min_count=1, workers=4, s
```

```
In [14]: ▶ def average_word_vectors(words, model, num_features):
    feature_vector = np.zeros((num_features,), dtype="float32")
    nwords = 0.

    for word in words:
        if word in model.wv:
            nwords += 1
            feature_vector = np.add(feature_vector, model.wv[word])

    if nwords:
        feature_vector = np.divide(feature_vector, nwords)

    return feature_vector
```

```
In [15]: ▶ X_train_cbow = [average_word_vectors(tokens, w2v_model_cbow, 100) for tokens in X_train]
    X_test_cbow = [average_word_vectors(tokens, w2v_model_cbow, 100) for tokens in X_test]
```

```
In [16]: ▶ label_encoder = LabelEncoder()
    y_train_encoded = label_encoder.fit_transform(y_train)
    y_test_encoded = label_encoder.transform(y_test)
```

```
In [17]: ▶ classifier_cbow = LogisticRegression(max_iter=1000)
    classifier_cbow.fit(X_train_cbow, y_train_encoded)
```

```
Out[17]: 

|                                   |                    |
|-----------------------------------|--------------------|
| ▼                                 | LogisticRegression |
| LogisticRegression(max_iter=1000) |                    |


```

```
In [18]: ▶ y_pred_cbow = classifier_cbow.predict(X_test_cbow)
    accuracy_cbow = accuracy_score(y_test_encoded, y_pred_cbow)
    print(f"Word2Vec CBOW Model Accuracy: {accuracy_cbow}")
```

Word2Vec CBOW Model Accuracy: 0.8365947145450877

```
In [19]: ▶ conf_matrix = confusion_matrix(y_test_encoded, y_pred_cbow)
    print("Confusion Matrix:")
    print(conf_matrix)
```

Confusion Matrix:

```
[[ 0 215  75]
 [ 0 3666 166]
 [ 0  354 481]]
```

```
In [20]: > class_report = classification_report(y_test_encoded, y_pred_cbow)
print("Classification Report:")
print(class_report)
```

```
Classification Report:
              precision    recall  f1-score   support

     0           0.00       0.00       0.00         290
     1           0.87       0.96       0.91        3832
     2           0.67       0.58       0.62         835

 accuracy                   0.84         4957
 macro avg           0.51       0.51       0.51         4957
 weighted avg        0.78       0.84       0.81         4957
```

```
C:\Users\chon2\AppData\Local\Programs\Python\Python39\lib\site-packages\sklearn\metrics\_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
C:\Users\chon2\AppData\Local\Programs\Python\Python39\lib\site-packages\sklearn\metrics\_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
C:\Users\chon2\AppData\Local\Programs\Python\Python39\lib\site-packages\sklearn\metrics\_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
```

## Word2Vect - SKIP GRAM

### STEPS:

- TRAINING THE WORD2VECT - CBOW MODEL
- AVERAGE WORD VECTORS FOR A SENTENCE
- TRANSFORMING TRAINING AND TEST DATA TO AVERAGE WORD VECTORS
- ENCODE LABELS TO NUMERIC VALUES
- CLASSIFICATION USING LOGISTIC REGRESSION
- EVALUATING THE MODEL
  - ACCURACY
  - CONFUSION MATRIX
  - CLASSIFICATION REPORT

```
In [21]: > from gensim.models import Word2Vec
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import LabelEncoder
import numpy as np
from sklearn.metrics import confusion_matrix, classification_report

from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

w2v_model_sg = Word2Vec(sentences=X_train, vector_size=100, window=5, min_count=1, workers=4, sg=
```

```
In [22]: ▶ def average_word_vectors(words, model, num_features):
    feature_vector = np.zeros((num_features,), dtype="float32")
    nwords = 0.

    for word in words:
        if word in model.wv:
            nwords += 1
            feature_vector = np.add(feature_vector, model.wv[word])

    if nwords:
        feature_vector = np.divide(feature_vector, nwords)

    return feature_vector
```

```
In [23]: ▶ X_train_sg = [average_word_vectors(tokens, w2v_model_sg, 100) for tokens in X_train]
    X_test_sg = [average_word_vectors(tokens, w2v_model_sg, 100) for tokens in X_test]
```

```
In [24]: ▶ label_encoder = LabelEncoder()
    y_train_encoded = label_encoder.fit_transform(y_train)
    y_test_encoded = label_encoder.transform(y_test)
```

```
In [25]: ▶ classifier_sg = LogisticRegression(max_iter=1000)
    classifier_sg.fit(X_train_sg, y_train_encoded)
```

```
Out[25]: LogisticRegression
LogisticRegression(max_iter=1000)
```

```
In [26]: ▶ y_pred_sg = classifier_sg.predict(X_test_sg)
    accuracy_sg = accuracy_score(y_test_encoded, y_pred_sg)
    print(f"Word2Vec Skip-gram Model Accuracy: {accuracy_sg}")
```

Word2Vec Skip-gram Model Accuracy: 0.8601977002219084

```
In [27]: ▶ conf_matrix_sg = confusion_matrix(y_test_encoded, y_pred_sg)
    print("Confusion Matrix:")
    print(conf_matrix_sg)
```

Confusion Matrix:

```
[[ 8 222 60]
 [ 19 3676 137]
 [ 4 251 580]]
```

```
In [28]: ▶ class_report_sg = classification_report(y_test_encoded, y_pred_sg)
    print("Classification Report:")
    print(class_report_sg)
```

Classification Report:

	precision	recall	f1-score	support
0	0.26	0.03	0.05	290
1	0.89	0.96	0.92	3832
2	0.75	0.69	0.72	835
accuracy			0.86	4957
macro avg	0.63	0.56	0.56	4957
weighted avg	0.83	0.86	0.84	4957

## FAST TEXT

### STEPS:

- LOADING THE FASTTEXT PRE-TRAINED MODEL
- FASTTEXT FOR A SENTENCE
- TRANSFORMING TRAINING AND TEST DATA TO FASTTEXT VECTORS
- ENCODE LABELS TO NUMERIC VALUES
- CLASSIFICATION USING LOGISTIC REGRESSION
- EVALUATING THE MODEL
  - ACCURACY
  - CONFUSION MATRIX
  - CLASSIFICATION REPORT

```
In [29]: from gensim.models import FastText
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import LabelEncoder
import numpy as np

from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

fasttext_model = FastText.load_fasttext_format('cc.en.300.bin')
```

C:\Users\chon2\AppData\Local\Temp\ipykernel\_2264\1509263375.py:12: DeprecationWarning: Call to deprecated `load\_fasttext\_format` (use load\_facebook\_vectors (to use pretrained embeddings) or load\_facebook\_model (to continue training with the loaded full model, more RAM) instead).

```
fasttext_model = FastText.load_fasttext_format('cc.en.300.bin')
```

```
In [30]: def get_fasttext_vectors(sentences, model):
vectors = []
for sentence in sentences:
    vector = np.mean([model.wv[word] for word in sentence if word in model.wv], axis=0)
    vectors.append(vector)
return vectors
```

```
In [31]: X_train_fasttext = get_fasttext_vectors(X_train, fasttext_model)
X_test_fasttext = get_fasttext_vectors(X_test, fasttext_model)
```

```
In [32]: label_encoder = LabelEncoder()
y_train_encoded = label_encoder.fit_transform(y_train)
y_test_encoded = label_encoder.transform(y_test)
```

```
In [33]: classifier_fasttext = LogisticRegression(max_iter=1000)
classifier_fasttext.fit(X_train_fasttext, y_train_encoded)
```

```
Out[33]: LogisticRegression
LogisticRegression(max_iter=1000)
```

```
In [34]: y_pred_fasttext = classifier_fasttext.predict(X_test_fasttext)
accuracy_fasttext = accuracy_score(y_test_encoded, y_pred_fasttext)
print(f"FastText Model Accuracy: {accuracy_fasttext}")
```

FastText Model Accuracy: 0.8793625176518055

```
In [35]: conf_matrix_fasttext = confusion_matrix(y_test_encoded, y_pred_fasttext)
print("Confusion Matrix:")
print(conf_matrix_fasttext)
```

Confusion Matrix:

```
[[ 33 223  34]
 [ 31 3698 103]
 [   7  200 628]]
```

```
In [36]: class_report_fasttext = classification_report(y_test_encoded, y_pred_fasttext)
print("Classification Report:")
print(class_report_fasttext)
```

```
Classification Report:
              precision    recall  f1-score   support

     0           0.46        0.11        0.18         290
     1           0.90        0.97        0.93        3832
     2           0.82        0.75        0.78         835

 accuracy              0.88         4957
 macro avg           0.73         0.61        0.63         4957
 weighted avg        0.86         0.88        0.86         4957
```

## CONVOLUTIONAL NEURAL NETWORK

### STEPS:

- TOKENIZE AND PAD SEQUENCES
- ENCODE LABELS TO NUMERIC VALUES
- BUILDING CNN MODEL
- EVALUATE AND PREDICT
- DECODE NUMERIC VALUES TO LABELS
- METRICS
  - ACCURACY
  - CONFUSION MATRIX
  - CLASSIFICATION REPORT

```
In [37]: from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Embedding, Conv1D, GlobalMaxPooling1D, Dense
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

tokenizer = Tokenizer()
tokenizer.fit_on_texts(X_train)

X_train_sequences = tokenizer.texts_to_sequences(X_train)
X_test_sequences = tokenizer.texts_to_sequences(X_test)

vocab_size = len(tokenizer.word_index) + 1

max_length = max(len(seq) for seq in X_train_sequences)

X_train_padded = pad_sequences(X_train_sequences, maxlen=max_length)
X_test_padded = pad_sequences(X_test_sequences, maxlen=max_length)
```

```
In [38]: label_encoder = LabelEncoder()
y_encoded = label_encoder.fit_transform(y)

X_train, X_test, y_train, y_test = train_test_split(X, y_encoded, test_size=0.2, random_state=42)
```



```

In [39]: embedding_dim = 100
         model = Sequential()
         model.add(Embedding(input_dim=vocab_size, output_dim=embedding_dim, input_length=max_length))
         model.add(Conv1D(filters=128, kernel_size=3, activation='relu'))
         model.add(GlobalMaxPooling1D())
         model.add(Dense(64, activation='relu'))

In [40]: model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])

In [41]: model.fit(X_train_padded, y_train_encoded, epochs=5, batch_size=32, validation_split=0.2)

Epoch 1/5
496/496 [=====] - 13s 25ms/step - loss: 3.9180 - accuracy: 0.1940 - val_loss: 4.1589 - val_accuracy: 0.0517
Epoch 2/5
496/496 [=====] - 11s 23ms/step - loss: 4.1589 - accuracy: 0.0590 - val_loss: 4.1589 - val_accuracy: 0.0517
Epoch 3/5
496/496 [=====] - 11s 23ms/step - loss: 4.1589 - accuracy: 0.0590 - val_loss: 4.1589 - val_accuracy: 0.0517
Epoch 4/5
496/496 [=====] - 11s 23ms/step - loss: 4.1589 - accuracy: 0.0590 - val_loss: 4.1589 - val_accuracy: 0.0517
Epoch 5/5
496/496 [=====] - 11s 23ms/step - loss: 4.1589 - accuracy: 0.0590 - val_loss: 4.1589 - val_accuracy: 0.0517

Out[41]: <keras.callbacks.History at 0x21fdd263ac0>

In [42]: loss, accuracy = model.evaluate(X_test_padded, y_test_encoded)

155/155 [=====] - 0s 2ms/step - loss: 4.1589 - accuracy: 0.0585

In [43]: y_pred_probs = model.predict(X_test_padded)
         y_pred = np.argmax(y_pred_probs, axis=1)

155/155 [=====] - 0s 2ms/step

In [44]: y_test_labels = label_encoder.inverse_transform(y_test_encoded)
         y_pred_labels = label_encoder.inverse_transform(y_pred)

In [45]: accuracy = accuracy_score(y_test_labels, y_pred_labels)
         print(f"Accuracy: {accuracy}")

Accuracy: 0.05850312689126488

In [46]: conf_matrix = confusion_matrix(y_test_labels, y_pred_labels)
         print("Confusion Matrix:")
         print(conf_matrix)

Confusion Matrix:
[[ 290   0   0]
 [3832   0   0]
 [ 835   0   0]]

```

```
In [47]: class_report = classification_report(y_test_labels, y_pred_labels)
print("Classification Report:")
print(class_report)
```

```
Classification Report:
              precision    recall  f1-score   support

     0       0.06        1.00      0.11         290
     1       0.00        0.00      0.00        3832
     2       0.00        0.00      0.00         835

 accuracy          0.06         4957
 macro avg       0.02         0.33         0.04         4957
 weighted avg    0.00         0.06         0.01         4957
```

```
C:\Users\chon2\AppData\Local\Programs\Python\Python39\lib\site-packages\sklearn\metrics\_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
C:\Users\chon2\AppData\Local\Programs\Python\Python39\lib\site-packages\sklearn\metrics\_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
C:\Users\chon2\AppData\Local\Programs\Python\Python39\lib\site-packages\sklearn\metrics\_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
```

## RECURRENT NEURAL NETWORK

### STEPS:

- TOKENIZE AND PAD SEQUENCES
- ENCODE LABELS TO NUMERIC VALUES
- BUILDING RNN MODEL
- EVALUATE AND PREDICT
- DECODE NUMERIC VALUES TO LABELS
- METRICS
  - ACCURACY
  - CONFUSION MATRIX
  - CLASSIFICATION REPORT

```
In [48]: from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Embedding, SimpleRNN, Dense
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences

from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

tokenizer = Tokenizer()
tokenizer.fit_on_texts(X_train)

X_train_sequences = tokenizer.texts_to_sequences(X_train)
X_test_sequences = tokenizer.texts_to_sequences(X_test)

vocab_size = len(tokenizer.word_index) + 1

max_length = max(len(seq) for seq in X_train_sequences)

X_train_padded = pad_sequences(X_train_sequences, maxlen=max_length)
X_test_padded = pad_sequences(X_test_sequences, maxlen=max_length)
```

```
In [49]: label_encoder = LabelEncoder()
y_encoded = label_encoder.fit_transform(y)

X_train, X_test, y_train, y_test = train_test_split(X, y_encoded, test_size=0.2, random_state=42)
```

```
In [50]: embedding_dim = 100
model = Sequential()
model.add(Embedding(input_dim=vocab_size, output_dim=embedding_dim, input_length=max_length))
model.add(SimpleRNN(units=64, activation='relu'))
model.add(Dense(64, activation='relu'))
```

```
In [51]: model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])
```

```
In [52]: model.fit(X_train_padded, y_train_encoded, epochs=5, batch_size=32, validation_split=0.2)

Epoch 1/5
496/496 [=====] - 13s 23ms/step - loss: 0.7662 - accuracy: 0.7639 - va
l_loss: 0.5860 - val_accuracy: 0.8058
Epoch 2/5
496/496 [=====] - 11s 23ms/step - loss: 0.4062 - accuracy: 0.8992 - va
l_loss: 0.5590 - val_accuracy: 0.8255
Epoch 3/5
496/496 [=====] - 11s 23ms/step - loss: 0.2375 - accuracy: 0.9444 - va
l_loss: 0.9007 - val_accuracy: 0.8719
Epoch 4/5
496/496 [=====] - 11s 23ms/step - loss: 0.1785 - accuracy: 0.9663 - va
l_loss: 1.4074 - val_accuracy: 0.8336
Epoch 5/5
496/496 [=====] - 11s 23ms/step - loss: 0.8098 - accuracy: 0.8436 - va
l_loss: 3.2249 - val_accuracy: 0.5436
```

```
Out[52]: <keras.callbacks.History at 0x21fa036df40>
```

```
In [53]: loss, accuracy = model.evaluate(X_test_padded, y_test_encoded)

155/155 [=====] - 0s 2ms/step - loss: 3.0184 - accuracy: 0.5568
```

```
In [54]: y_pred_probs = model.predict(X_test_padded)
y_pred = np.argmax(y_pred_probs, axis=1)

155/155 [=====] - 0s 2ms/step
```

```
In [55]: y_test_labels = label_encoder.inverse_transform(y_test_encoded)
y_pred_labels = label_encoder.inverse_transform(y_pred)
```

```
In [56]: accuracy = accuracy_score(y_test_labels, y_pred_labels)
print(f"Accuracy: {accuracy}")
```

Accuracy: 0.5567883800685899

```
In [57]: conf_matrix = confusion_matrix(y_test_labels, y_pred_labels)
print("Confusion Matrix:")
print(conf_matrix)
```

Confusion Matrix:  
[[ 145 76 69]  
 [1451 1956 425]  
 [ 113 63 659]]

```
In [58]: class_report = classification_report(y_test_labels, y_pred_labels)
print("Classification Report:")
print(class_report)
```

Classification Report:

	precision	recall	f1-score	support
0	0.08	0.50	0.15	290
1	0.93	0.51	0.66	3832
2	0.57	0.79	0.66	835
accuracy			0.56	4957
macro avg	0.53	0.60	0.49	4957
weighted avg	0.82	0.56	0.63	4957